

1. kubectl version --short

```
> kubectl version --short  
Client Version: v1.22.1  
Server Version: v1.21.2-eks-0389ca3
```

This command, helps us see **which version of the API server** is running.

This gives us important information when we're troubleshooting specific errors, and it's very useful to know if we're on an older cluster like 1.16.

2. kubectl cluster-info

```
> kubectl cluster-info
Kubernetes control plane is running at https://20440DACB4861717B9FEBAB3850935B0.gr7.us-east-2.eks.amazonaws.com
CoreDNS is running at https://20440DACB4861717B9FEBAB3850935B0.gr7.us-east-2.eks.amazonaws.com/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

Next, we should understand where the cluster is running and if CoreDNS is running.

You can parse the control plane URL to know if you're dealing with a hosted cluster or something on-premises.

3. kubectl get componentstatus

```
> kubectl get componentstatus
Warning: v1 ComponentStatus is deprecated in v1.19+
NAME                                STATUS    MESSAGE                                           ERROR
scheduler                           Healthy   ok
controller-manager                  Healthy   ok
etcd-0                              Healthy   {"health":"true"}
```

This command will be the easiest way to discover if your **scheduler**, **controller-manager** and **etcd node(s)** are **healthy**. These are all critical control plane components to run your pods. Look for errors on the components that don't show an "ok" status.

An alternative option to see other health endpoints, including etcd, is

kubectl get --raw '/healthz?verbose':

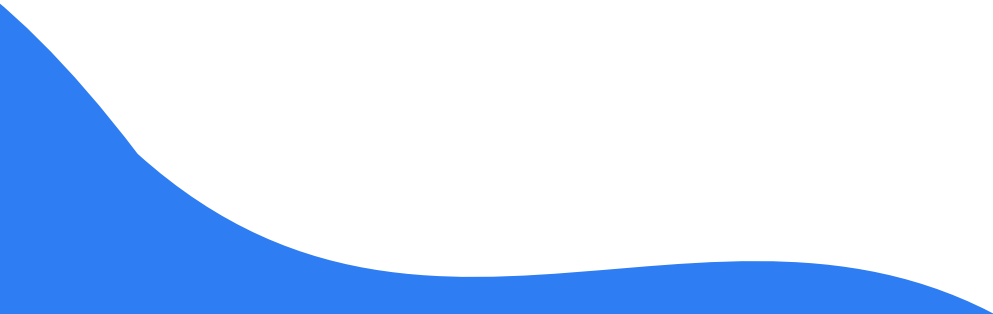

```
> kubectl get --raw '/healthz?verbose'
[+]ping ok
[+]log ok
[+]etcd ok
[+]poststarthook/start-kube-apiserver-admission-initializer ok
[+]poststarthook/generic-apiserver-start-informers ok
[+]poststarthook/priority-and-fairness-config-consumer ok
[+]poststarthook/priority-and-fairness-filter ok
[+]poststarthook/start-apiextensions-informers ok
[+]poststarthook/start-apiextensions-controllers ok
[+]poststarthook/crd-informer-synced ok
[+]poststarthook/bootstrap-controller ok
[+]poststarthook/rbac/bootstrap-roles ok
[+]poststarthook/scheduling/bootstrap-system-priority-classes ok
[+]poststarthook/priority-and-fairness-config-producer ok
[+]poststarthook/start-cluster-authentication-info-controller ok
[+]poststarthook/aggregator-reload-proxy-client-cert ok
[+]poststarthook/start-kube-aggregator-informers ok
[+]poststarthook/apiservice-registration-controller ok
[+]poststarthook/apiservice-status-available-controller ok
[+]poststarthook/kube-apiserver-autoregistration ok
[+]autoregister-completion ok
[+]poststarthook/apiservice-openapi-controller ok
healthz check passed
```

4. kubectl api-resources -o wide --sort-by name

```
> kubectl api-resources -o wide --sort-by name
```

NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND	VERBS
alertmanagerconfigs		monitoring.coreos.com/v1alpha1	true	AlertmanagerConfig	[delete deletetecollection get list patch create update watch]
alertmanagers		monitoring.coreos.com/v1	true	Alertmanager	[delete deletetecollection get list patch create update watch]
apiservices		apiregistration.k8s.io/v1	false	APIService	[create delete deletetecollection get list patch update watch]
bindings		v1	true	Binding	[create]
certificatesigningrequests	csr	certificates.k8s.io/v1	false	CertificateSigningRequest	[create delete deletetecollection get list patch update watch]
clusterrolebindings		rbac.authorization.k8s.io/v1	false	ClusterRoleBinding	[create delete deletetecollection get list patch update watch]
clusterroles		rbac.authorization.k8s.io/v1	false	ClusterRole	[create delete deletetecollection get list patch update watch]
componentstatuses	cs	v1	false	ComponentStatus	[get list]
configmaps	cm	v1	true	ConfigMap	[create delete deletetecollection get list patch update watch]
controllerrevisions		apps/v1	true	ControllerRevision	[create delete deletetecollection get list patch update watch]
cronjobs	cj	batch/v1	true	CronJob	[create delete deletetecollection get list patch update watch]
csidrivers		storage.k8s.io/v1	false	CSIDriver	[create delete deletetecollection get list patch update watch]
csinodes		storage.k8s.io/v1	false	CSINode	[create delete deletetecollection get list patch update watch]
csistoragecapacities		storage.k8s.io/v1beta1	true	CSIStorageCapacity	[create delete deletetecollection get list patch update watch]
customresourcedefinitions	crd, crds	apiextensions.k8s.io/v1	false	CustomResourceDefinition	[create delete deletetecollection get list patch update watch]
daemonsets	ds	apps/v1	true	DaemonSet	[create delete deletetecollection get list patch update watch]
deployments	deploy	apps/v1	true	Deployment	[create delete deletetecollection get list patch update watch]
endpoints	ep	v1	true	Endpoints	[create delete deletetecollection get list patch update watch]
endpointslices		discovery.k8s.io/v1	true	EndpointSlice	[create delete deletetecollection get list patch update watch]
eniconfigs		crd.k8s.amazonaws.com/v1alpha1	false	ENIConfig	[delete deletetecollection get list patch create update watch]
events	ev	v1	true	Event	[create delete deletetecollection get list patch update watch]
events	ev	events.k8s.io/v1	true	Event	[create delete deletetecollection get list patch update watch]
flowschemas		flowcontrol.apiserver.k8s.io/v1beta1	false	FlowSchema	[create delete deletetecollection get list patch update watch]
horizontalpodautoscalers	hpa	autoscaling/v1	true	HorizontalPodAutoscaler	[create delete deletetecollection get list patch update watch]
ingressclasses		networking.k8s.io/v1	false	IngressClass	[create delete deletetecollection get list patch update watch]
ingresses	ing	networking.k8s.io/v1	true	Ingress	[create delete deletetecollection get list patch update watch]
ingresses	ing	extensions/v1beta1	true	Ingress	[create delete deletetecollection get list patch update watch]
jobs		batch/v1	true	Job	[create delete deletetecollection get list patch update watch]
leases		coordination.k8s.io/v1	true	Lease	[create delete deletetecollection get list patch update watch]
limitranges	limits	v1	true	LimitRange	[create delete deletetecollection get list patch update watch]
localsubjectaccessreviews		authorization.k8s.io/v1	true	LocalSubjectAccessReview	[create]
mutatingwebhookconfigurations		admissionregistration.k8s.io/v1	false	MutatingWebhookConfiguration	[create delete deletetecollection get list patch update watch]
namespaces	ns	v1	false	Namespace	[create delete get list patch update watch]

I like to list all the resources sorted by name for consistency. It's easier for me to scan the resources in alphabetical order. Adding **-o wide** will show the verbs available on each resource.



Using this command will tell you what CRDs (custom resource definitions) have been installed in your cluster and what API version each resource is at.

5. kubectl get events -A

NAMESPACE	LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
default	48m	Normal	RegisteredNode	node/ip-192-168-103-110.us-east-2.compute.internal	Node ip-192-168-103-110.us-east-2.compute.internal event: Registered N
ode ip-192-168-103-110.us-east-2.compute.internal in Controller					
default	47m	Normal	Starting	node/ip-192-168-103-110.us-east-2.compute.internal	Starting kubelet.
default	47m	Warning	InvalidDiskCapacity	node/ip-192-168-103-110.us-east-2.compute.internal	invalid capacity 0 on image filesystem
default	47m	Normal	NodeHasSufficientMemory	node/ip-192-168-103-110.us-east-2.compute.internal	Node ip-192-168-103-110.us-east-2.compute.internal status is now: Node
HasSufficientMemory					
default	47m	Normal	NodeHasNoDiskPressure	node/ip-192-168-103-110.us-east-2.compute.internal	Node ip-192-168-103-110.us-east-2.compute.internal status is now: Node
HasNoDiskPressure					
default	47m	Normal	NodeHasSufficientPID	node/ip-192-168-103-110.us-east-2.compute.internal	Node ip-192-168-103-110.us-east-2.compute.internal status is now: Node
HasSufficientPID					
default	47m	Normal	NodeAllocatableEnforced	node/ip-192-168-103-110.us-east-2.compute.internal	Updated Node Allocatable limit across pods
default	47m	Normal	NodeNotReady	node/ip-192-168-103-110.us-east-2.compute.internal	Node ip-192-168-103-110.us-east-2.compute.internal status is now: Node
NotReady					
default	47m	Normal	Starting	node/ip-192-168-103-110.us-east-2.compute.internal	Starting kube-proxy.
default	47m	Normal	NodeReady	node/ip-192-168-103-110.us-east-2.compute.internal	Node ip-192-168-103-110.us-east-2.compute.internal status is now: Node
Ready					
default	38m	Normal	NodeNotSchedulable	node/ip-192-168-103-110.us-east-2.compute.internal	Node ip-192-168-103-110.us-east-2.compute.internal status is now: Node
NotSchedulable					
default	38m	Normal	RemovingNode	node/ip-192-168-103-110.us-east-2.compute.internal	Node ip-192-168-103-110.us-east-2.compute.internal event: Removing Nod
e ip-192-168-103-110.us-east-2.compute.internal from Controller					
default	28m	Normal	RegisteredNode	node/ip-192-168-110-139.us-east-2.compute.internal	Node ip-192-168-110-139.us-east-2.compute.internal event: Registered N
ode ip-192-168-110-139.us-east-2.compute.internal in Controller					
default	27m	Normal	Starting	node/ip-192-168-110-139.us-east-2.compute.internal	Starting kubelet.
default	27m	Warning	InvalidDiskCapacity	node/ip-192-168-110-139.us-east-2.compute.internal	invalid capacity 0 on image filesystem
default	27m	Normal	NodeHasSufficientMemory	node/ip-192-168-110-139.us-east-2.compute.internal	Node ip-192-168-110-139.us-east-2.compute.internal status is now: Node
HasSufficientMemory					
default	27m	Normal	NodeHasNoDiskPressure	node/ip-192-168-110-139.us-east-2.compute.internal	Node ip-192-168-110-139.us-east-2.compute.internal status is now: Node
HasNoDiskPressure					
default	27m	Normal	NodeHasSufficientPID	node/ip-192-168-110-139.us-east-2.compute.internal	Node ip-192-168-110-139.us-east-2.compute.internal status is now: Node
HasSufficientPID					
default	27m	Normal	NodeAllocatableEnforced	node/ip-192-168-110-139.us-east-2.compute.internal	Updated Node Allocatable limit across pods
default	27m	Normal	NodeNotReady	node/ip-192-168-110-139.us-east-2.compute.internal	Node ip-192-168-110-139.us-east-2.compute.internal status is now: Node
NotReady					
default	27m	Normal	Starting	node/ip-192-168-110-139.us-east-2.compute.internal	Starting kube-proxy.
default	27m	Normal	NodeReady	node/ip-192-168-110-139.us-east-2.compute.internal	Node ip-192-168-110-139.us-east-2.compute.internal status is now: Node
Ready					

Now that we have an idea of what's running in the cluster, we should look at what's happening. If something broke recently, **you can look at the cluster events** to see what was happening before and after things broke.

6. kubectl get nodes -o wide

```
> kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME
fargate-ip-192-168-107-143.us-east-2.compute.internal	Ready	<none>	5d7h	v1.21.2-eks-55daa9d	192.168.107.143	<none>	Amazon Linux 2	4.14.243-185.433.amzn2.x86_64	containerd://1.4.6
fargate-ip-192-168-124-110.us-east-2.compute.internal	Ready	<none>	5d7h	v1.21.2-eks-55daa9d	192.168.124.110	<none>	Amazon Linux 2	4.14.243-185.433.amzn2.x86_64	containerd://1.4.6
fargate-ip-192-168-133-244.us-east-2.compute.internal	Ready	<none>	5d7h	v1.21.2-eks-55daa9d	192.168.133.244	<none>	Amazon Linux 2	4.14.243-185.433.amzn2.x86_64	containerd://1.4.6
fargate-ip-192-168-149-111.us-east-2.compute.internal	Ready	<none>	5d7h	v1.21.2-eks-55daa9d	192.168.149.111	<none>	Amazon Linux 2	4.14.243-185.433.amzn2.x86_64	containerd://1.4.6
fargate-ip-192-168-150-136.us-east-2.compute.internal	Ready	<none>	5d7h	v1.21.2-eks-55daa9d	192.168.150.136	<none>	Amazon Linux 2	4.14.243-185.433.amzn2.x86_64	containerd://1.4.6
ip-192-168-135-9.us-east-2.compute.internal	Ready, SchedulingDisabled	<none>	68d	v1.21.2	192.168.135.9	<none>	Bottlerocket OS 1.1.4	5.10.50	containerd://1.4.8+bottlerocket
ip-192-168-137-131.us-east-2.compute.internal	NotReady, SchedulingDisabled	<none>	43d	<none>	<none>	<none>	<unknown>	<unknown>	<unknown>
ip-192-168-152-72.us-east-2.compute.internal	NotReady, SchedulingDisabled	<none>	72d	v1.21.2	192.168.152.72	<none>	Bottlerocket OS 1.1.4	5.10.50	containerd://1.4.8+bottlerocket
ip-192-168-166-136.us-east-2.compute.internal	Ready, SchedulingDisabled	<none>	72d	v1.21.5	192.168.166.136	<none>	Bottlerocket OS 1.2.1	5.10.50	containerd://1.4.8+bottlerocket
ip-192-168-168-18.us-east-2.compute.internal	NotReady, SchedulingDisabled	<none>	68d	<none>	<none>	<none>	<unknown>	<unknown>	<unknown>
ip-192-168-176-99.us-east-2.compute.internal	NotReady, SchedulingDisabled	<none>	45d	<none>	<none>	<none>	<unknown>	<unknown>	<unknown>
ip-192-168-186-164.us-east-2.compute.internal	Ready	<none>	93s	v1.21.5	192.168.186.164	<none>	Bottlerocket OS 1.2.1	5.10.50	containerd://1.4.8+bottlerocket

Nodes are a first-class resource inside Kubernetes and are fundamental for pods to run.

Using the **-o wide** option will tell us additional details like operating system (OS), IP address, and container runtime.

7. kubectl get pods -A -o wide

```
> kubectl get po -A -o wide
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
default	alertmanager-prom-kube-prometheus-stack-alertmanager-0	0/2	Terminating	0	68d	<none>	ip-192-168-168-18.us-east-2.compute.internal	<none>	<none>
default	kube-ops-view-5b5d9b6bf8-4l8hj	1/1	Terminating	0	10m	192.168.175.85	ip-192-168-186-164.us-east-2.compute.internal	<none>	<none>
default	kube-ops-view-5b5d9b6bf8-gmwfk	0/1	Terminating	0	45d	<none>	ip-192-168-176-99.us-east-2.compute.internal	<none>	<none>
default	kube-ops-view-5b5d9b6bf8-kjwbj	0/1	Terminating	0	43d	<none>	ip-192-168-137-131.us-east-2.compute.internal	<none>	<none>
default	kube-ops-view-5b5d9b6bf8-m8zhr	0/1	Pending	0	4s	<none>	<none>	<none>	<none>
default	prom-grafana-5964dd8cb6-9gq28	2/2	Running	0	68d	192.168.149.238	ip-192-168-135-9.us-east-2.compute.internal	<none>	<none>
default	prom-grafana-5964dd8cb6-mdkfv	2/2	Terminating	0	72d	192.168.143.221	ip-192-168-152-72.us-east-2.compute.internal	<none>	<none>
default	prom-kube-prometheus-stack-operator-fb7c484b9-9qp67	0/1	Terminating	0	43d	<none>	ip-192-168-137-131.us-east-2.compute.internal	<none>	<none>
default	prom-kube-prometheus-stack-operator-fb7c484b9-lgskh	0/1	Pending	0	4s	<none>	<none>	<none>	<none>
default	prom-kube-prometheus-stack-operator-fb7c484b9-zqwn9	0/1	Terminating	0	45d	<none>	ip-192-168-176-99.us-east-2.compute.internal	<none>	<none>
default	prom-kube-state-metrics-695c5f66cc-26dq2	0/1	Terminating	0	10m	192.168.179.191	ip-192-168-186-164.us-east-2.compute.internal	<none>	<none>
default	prom-kube-state-metrics-695c5f66cc-7bt7h	0/1	Terminating	0	43d	<none>	ip-192-168-137-131.us-east-2.compute.internal	<none>	<none>
default	prom-kube-state-metrics-695c5f66cc-fgfnx	0/1	Terminating	0	45d	<none>	ip-192-168-176-99.us-east-2.compute.internal	<none>	<none>
default	prom-kube-state-metrics-695c5f66cc-q52js	0/1	Pending	0	4s	<none>	<none>	<none>	<none>
default	prom-prometheus-node-exporter-2dzww	1/1	Running	0	68d	192.168.135.9	ip-192-168-135-9.us-east-2.compute.internal	<none>	<none>
default	prom-prometheus-node-exporter-4hn8n	1/1	Running	0	72d	192.168.166.136	ip-192-168-166-136.us-east-2.compute.internal	<none>	<none>
default	prom-prometheus-node-exporter-4lhtt	0/1	Pending	0	5d7h	<none>	<none>	<none>	<none>
default	prom-prometheus-node-exporter-6rpgg	0/1	Pending	0	43d	<none>	ip-192-168-137-131.us-east-2.compute.internal	<none>	<none>
default	prom-prometheus-node-exporter-gjft7	0/1	Pending	0	5d7h	<none>	<none>	<none>	<none>
default	prom-prometheus-node-exporter-k9nxd	0/1	Pending	0	5d7h	<none>	<none>	<none>	<none>
default	prom-prometheus-node-exporter-kfvjn	1/1	Running	0	72d	192.168.152.72	ip-192-168-152-72.us-east-2.compute.internal	<none>	<none>
default	prom-prometheus-node-exporter-lqbbb	1/1	Running	0	10m	192.168.186.164	ip-192-168-186-164.us-east-2.compute.internal	<none>	<none>
default	prom-prometheus-node-exporter-pp5qq	0/1	Pending	0	68d	<none>	ip-192-168-168-18.us-east-2.compute.internal	<none>	<none>
default	prom-prometheus-node-exporter-qsnfd	0/1	Pending	0	5d7h	<none>	<none>	<none>	<none>
default	prom-prometheus-node-exporter-xpv4w	0/1	Pending	0	45d	<none>	ip-192-168-176-99.us-east-2.compute.internal	<none>	<none>
default	prom-prometheus-node-exporter-zgb67	0/1	Pending	0	5d7h	<none>	<none>	<none>	<none>
default	prometheus-prom-kube-prometheus-stack-prometheus-0	0/2	Terminating	0	68d	<none>	ip-192-168-168-18.us-east-2.compute.internal	<none>	<none>
karpenter	karpenter-controller-756fdd7447-7qn9f	1/1	Running	0	5d7h	192.168.107.143	fargate-ip-192-168-107-143.us-east-2.compute.internal	<none>	<none>
karpenter	karpenter-webhook-67f4fb4dd9-pwvtt	1/1	Running	0	5d7h	192.168.150.136	fargate-ip-192-168-150-136.us-east-2.compute.internal	<none>	<none>
kube-system	aws-node-57b9l	1/1	Running	0	9m19s	192.168.186.164	ip-192-168-186-164.us-east-2.compute.internal	<none>	<none>
kube-system	aws-node-jlcbp	1/1	Running	0	5d8h	192.168.166.136	ip-192-168-166-136.us-east-2.compute.internal	<none>	<none>
kube-system	aws-node-lhc85	1/1	Running	0	72d	192.168.152.72	ip-192-168-152-72.us-east-2.compute.internal	<none>	<none>
kube-system	aws-node-nncs8	1/1	Running	0	68d	192.168.135.9	ip-192-168-135-9.us-east-2.compute.internal	<none>	<none>
kube-system	coredns-697445b7b9-cf6ft	1/1	Running	0	5d7h	192.168.124.110	fargate-ip-192-168-124-110.us-east-2.compute.internal	<none>	<none>
kube-system	coredns-697445b7b9-k7kv2	1/1	Running	0	5d7h	192.168.133.244	fargate-ip-192-168-133-244.us-east-2.compute.internal	<none>	<none>
kube-system	kube-proxy-btqlz	1/1	Running	0	68d	192.168.135.9	ip-192-168-135-9.us-east-2.compute.internal	<none>	<none>
kube-system	kube-proxy-c25x8	1/1	Running	0	72d	192.168.152.72	ip-192-168-152-72.us-east-2.compute.internal	<none>	<none>
kube-system	kube-proxy-df8bk	1/1	Running	0	5d8h	192.168.166.136	ip-192-168-166-136.us-east-2.compute.internal	<none>	<none>
kube-system	kube-proxy-jxvb4	1/1	Running	0	9m19s	192.168.186.164	ip-192-168-186-164.us-east-2.compute.internal	<none>	<none>
kube-system	metrics-server-6dfddc5fb8-mtqvw	1/1	Running	0	5d7h	192.168.149.111	fargate-ip-192-168-149-111.us-east-2.compute.internal	<none>	<none>

Using -A will list pods in all namespaces and -o wide will show us IP addresses, nodes, and where the pods are nominated. Using the information from listing nodes, you can look at which pods are failing on which nodes.

8. `kubectl run a --image alpine --command -- /bin/sleep 1d`

```
> kubectl get po a
```

NAME	READY	STATUS	RESTARTS	AGE
a	1/1	Running	0	6s

Sometimes, the best way you can debug something is to start with the simplest example. This command doesn't have any direct output, **but you should see a running pod named "a" from it.**

```
Events:
  Type    Reason      Age    From          Message
  ----    -
  Normal  Scheduled   13s    default-scheduler Successfully assigned default/alpine to ip-192-168-173-156.us-east-2.compute.internal
  Normal  Pulled      12s    kubelet       Successfully pulled image "alpine" in 204.535296ms
  Normal  Pulling     11s (x2 over 12s)  kubelet       Pulling image "alpine"
  Normal  Created     11s (x2 over 12s)  kubelet       Created container alpine
  Warning Failed      11s (x2 over 12s)  kubelet       Error: failed to create containerd task: OCI runtime create failed: container_linux.go:380: starting container
process caused: exec: "/bin/sleep 1d": stat /bin/sleep 1d: no such file or directory: unknown
  Normal  Pulled      11s    kubelet       Successfully pulled image "alpine" in 221.172265ms
  Warning BackOff     9s (x2 over 10s)  kubelet       Back-off restarting failed container
```

If for some reason you don't see a running pod from this command, then using **kubectl describe po a** is your next-best option.

Look at the events to find errors for what might have gone wrong.

Credit: Justin Garrison, The new stack