



{K}ODE{K}LOUD

www.kodecloud.com

cks.kodecloud.com

Disclaimer

THE INFORMATION FOUND ON THE WEBSITE, E-LEARNING PLATFORM AND WITHIN THE ONLINE COURSES ARE FOR INFORMATIONAL PURPOSES ONLY. KODEKLOUD WILL NOT BE HELD RESPONSIBLE FOR ANY DAMAGES THAT MAY BE INCURRED BY YOU AS A RESULT OF YOUR USE OF SUCH INFORMATION. ALL INFORMATION AND CONTENT ON THE WEBSITE, E-LEARNING PLATFORM AND ONLINE COURSE IS COPYRIGHTED, AND MAY NOT BE REPUBLISHED, COPIED, SOLD OR POSTED ANYWHERE ONLINE OR IN PRINT. KODEKLOUD RESERVES THE RIGHT TO TAKE THE NECESSARY LEGAL ACTION TO PREVENT YOU FROM (RE)-PUBLISHING, COPYING, SELLING, POSTING OR PRINTING ANY COPYRIGHTED INFORMATION AND CONTENT AVAILABLE ON THE WEBSITE, E-LEARNING PLATFORM AND ONLINE COURSE.

For the full terms & conditions visit terms.kodekloud.com

For questions write to support@kodekloud.com

Notice

- This presentation is to refer to course contents only.
- Some of the slides are meant to be animated. So may not be displayed correctly.
- Do not copy and paste command, code or YAML files from this file as it may not be in the right format and may contain hidden characters
- For code refer to the solutions in the lab or the Git repository associated with this course or official Kubernetes documentation pages.
- Some of the code in this deck maybe hidden for brevity

<https://github.com/kodekloudhub/certified-kubernetes-security-specialist-cks-course>



Perform Behaviour Analytics of syscalls



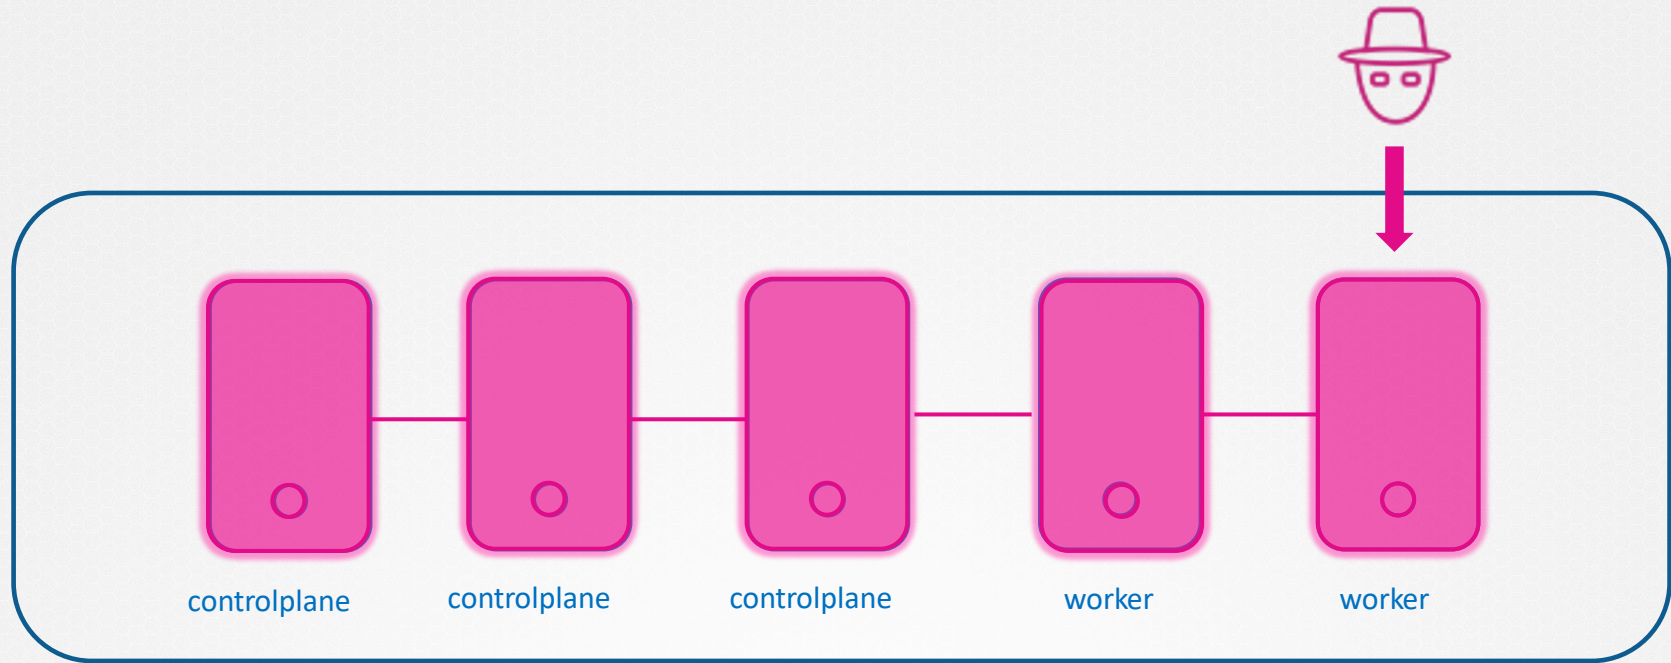
Securing Cluster

Sandboxing Techniques

Restricting Network Access

Minimizing Microservices
Vulnerability

MTLS Encryption





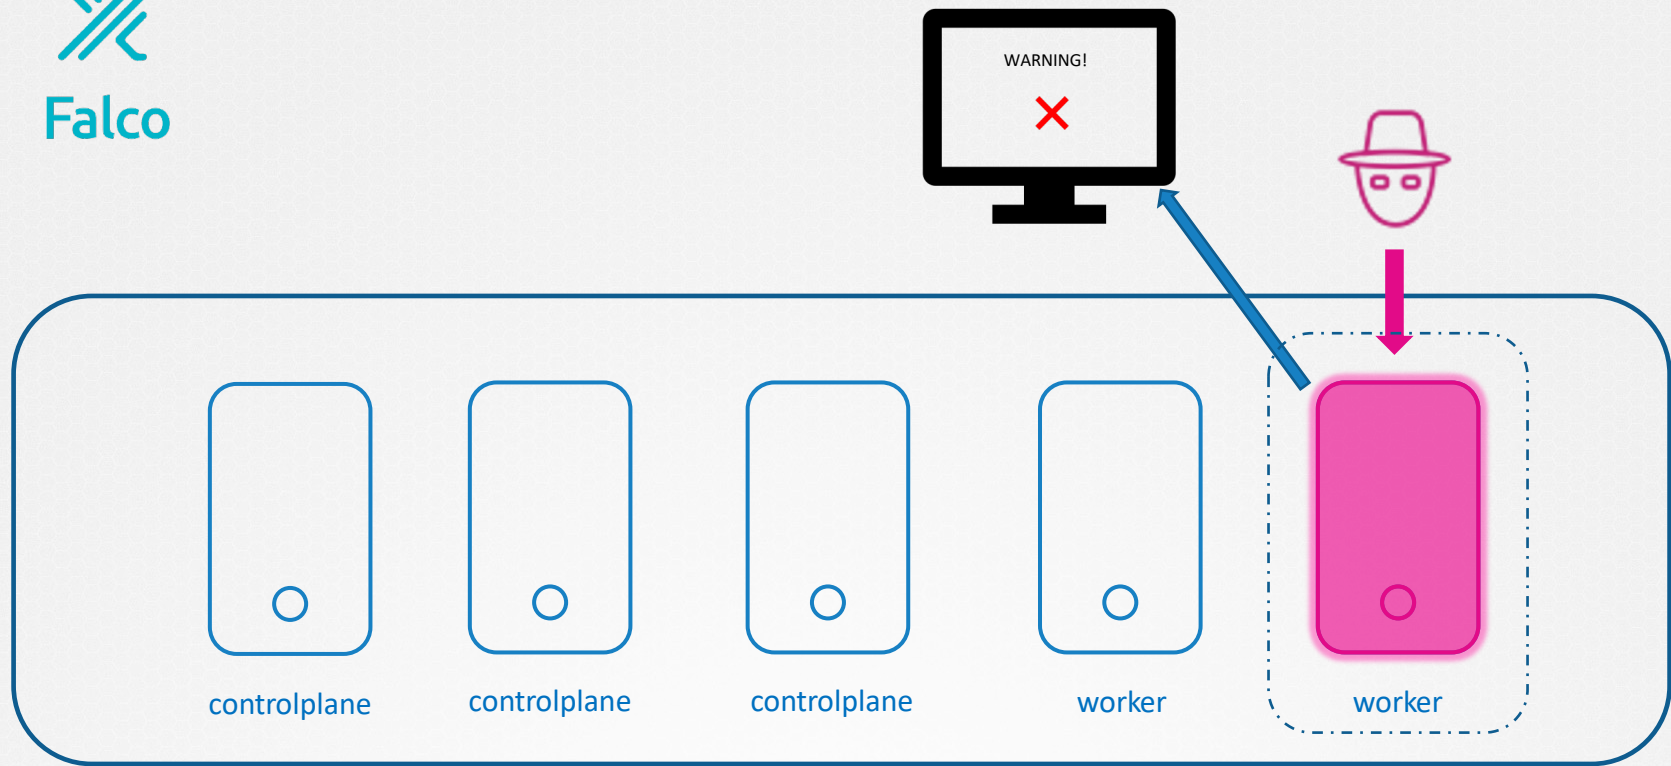
Instant Notifications



Revert Transactions

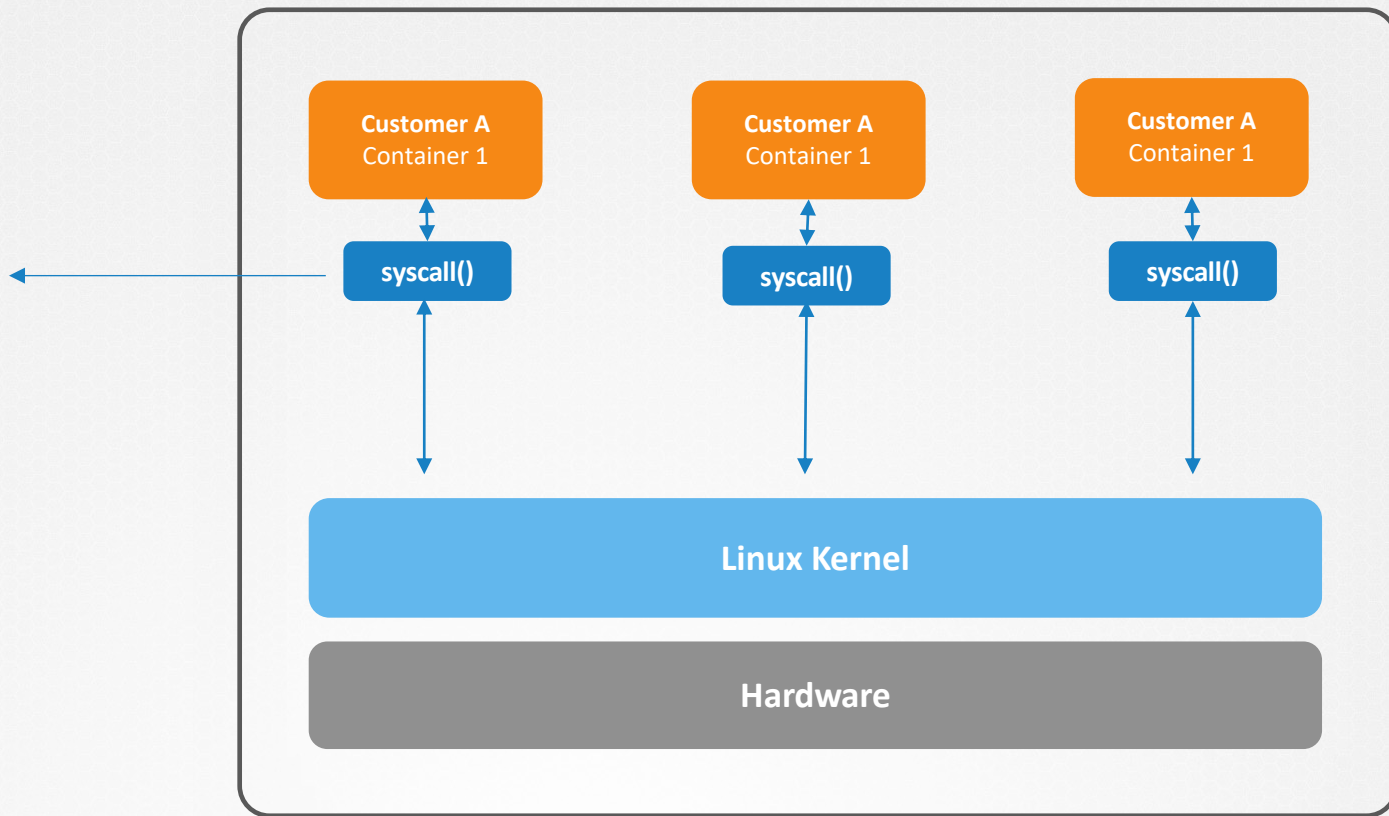



Transaction Limits





SYSCALL NAME
close
nanosleep
fcntl
fstatfs
getdents64
exit_group
epoll_ctl
openat





```
▶ kubectl exec -ti nginx-master -- bash  
# cat /etc/shadow
```



```
▶ > /opt/logs/audit.log
```



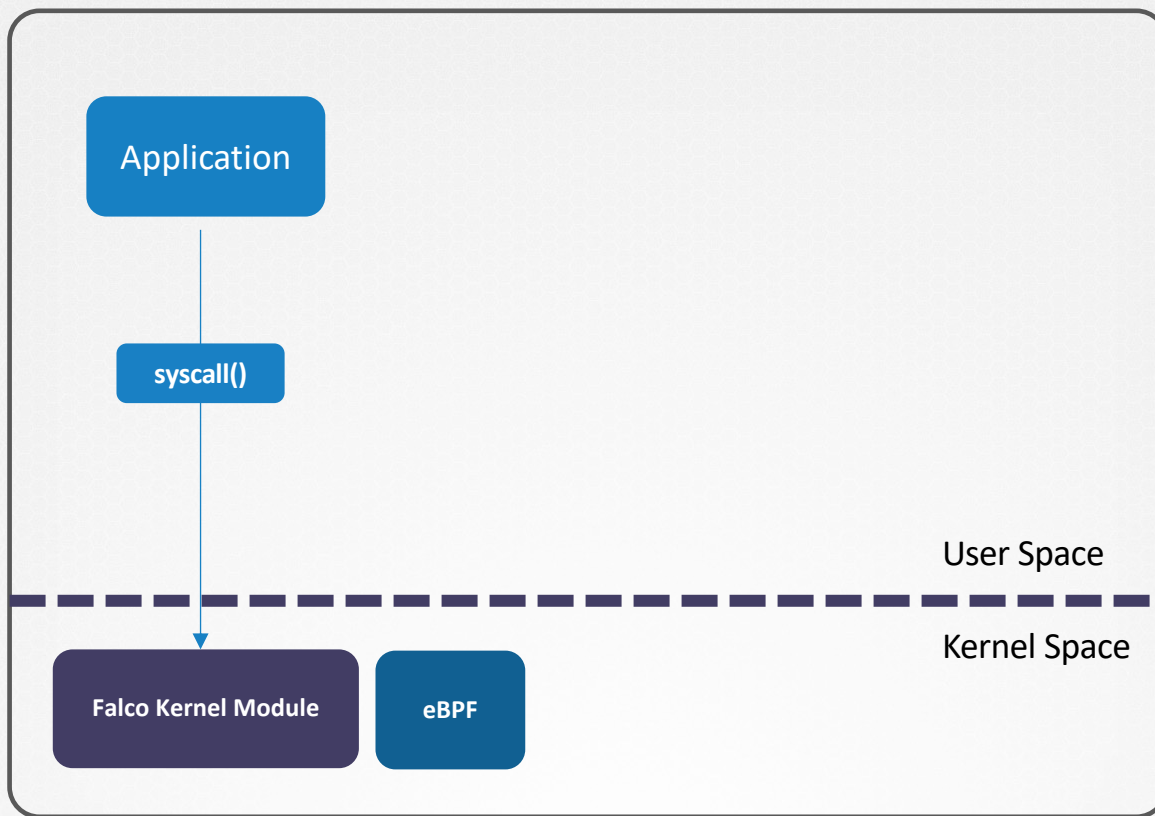
{K}ODE{K}LOUD

www.kodecloud.com

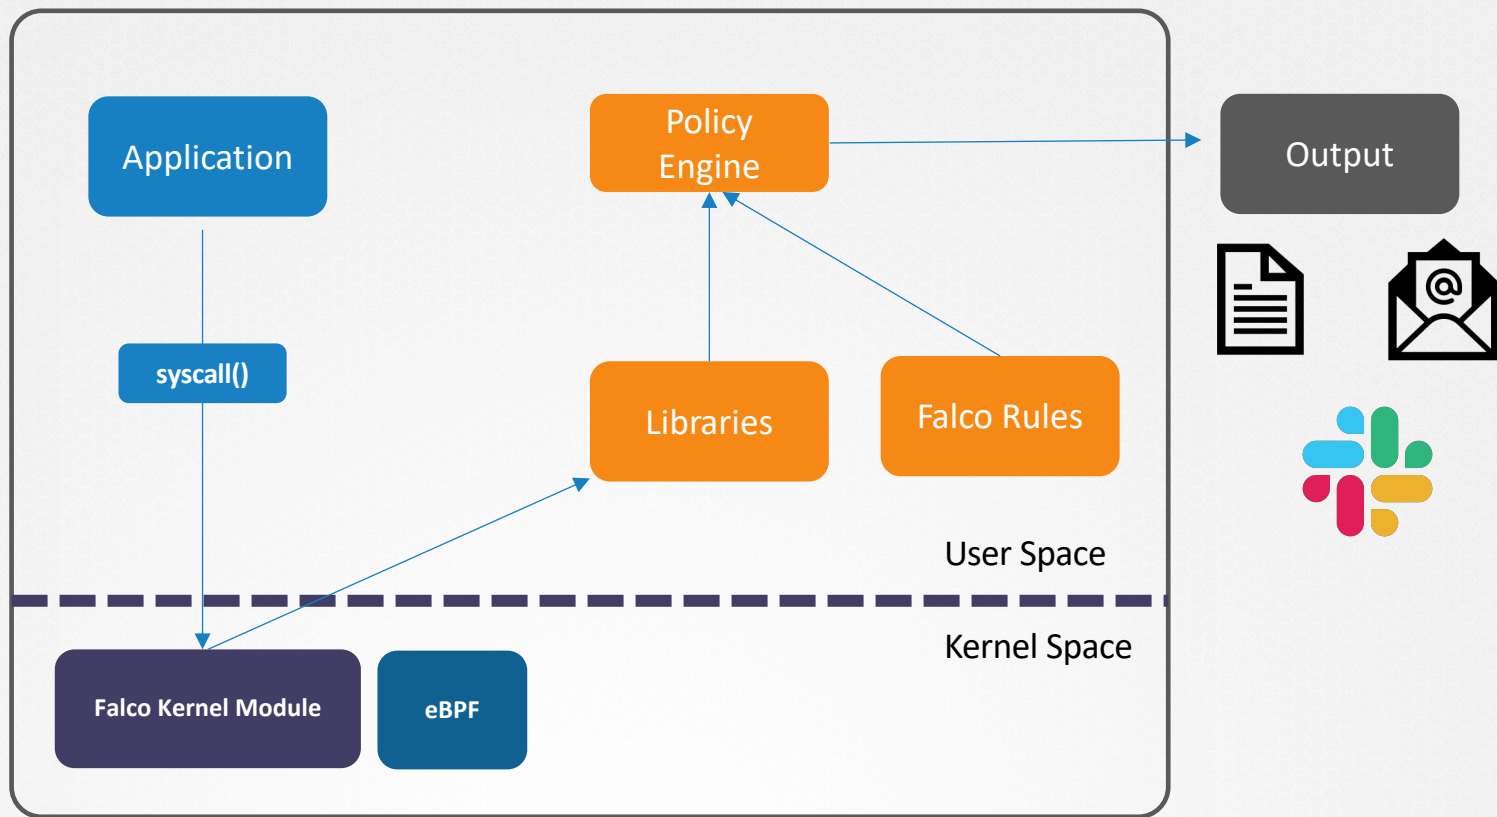
Falco Overview and Installation



Falco Architecture



Falco Architecture



Install as a Package

```
▶ curl -s https://falco.org/repo/falcosecurity-3672BA8F.asc | apt-key add -
```

```
▶ echo "deb https://download.falco.org/packages/deb stable main" | tee -a /etc/apt/sources.list.d/falcosecurity.list
```

```
▶ apt update -y
```

```
▶ apt get install -y linux-headers-$(uname -r)
```

```
▶ apt install -y falco
```

```
▶ systemctl start falco
```

Install as a DaemonSet

```
▶ helm repo add falcosecurity https://falcosecurity.github.io/charts
```

```
▶ helm repo update
```

```
▶ helm install falco falcosecurity/falco
```

NAME: falco

LAST DEPLOYED: Wed Mar 7 20:19:25 2021

NAMESPACE: default

STATUS: deployed

REVISION: 1

TEST SUITE: None

NOTES:

Falco agents are spinning up on each node in your cluster. After a few seconds, they are going to start monitoring your containers looking for security issues.

No further action should be required.

I Install as a DaemonSet

```
▶ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
falco-7grdt	1/1	Running	0	2m21s
falco-tm928	1/1	Running	0	2m21s




{K}ODE{K}LOUD

www.kodecloud.com

Use Falco to Detect Threats



node01

 `systemctl status falco`

- `falco.service` - Falco: Container Native Runtime Security
Loaded: loaded (/usr/lib/systemd/system/falco.service; enabled; vendor preset: enabled)
Active: active (running) since Tue 2021-04-13 20:42:45 UTC; 1min 2s ago
Docs: <https://falco.org/docs/>
Process: 17981 ExecStartPre=/sbin/modprobe falco (code=exited, status=0/SUCCESS)
Main PID: 17994 (falco)
Tasks: 6 (limit: 4678)
CGroup: /system.slice/falco.service
└─17994 /usr/bin/falco --pidfile=/var/run/falco.pid -c /etc/falco/falco.yaml

 `kubectl run nginx --image=nginx``pod/nginx created` `kubectl get pods -o wide`

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
nginx	1/1	Running	0	6m1s	10.244.1.3	node01	<none>	<none>

node01

```
journalctl -fu falco
```

```
.  
.
.
```

```
22:57:09.163982780: Notice A shell was spawned in a container with an attached terminal (user=root
user_loginuid=-1 k8s.ns=default k8s.pod=nginx container=c73d9fc1a75d shell=bash parent=runc
cmdline=bash terminal=34816 container_id=c73d9fc1a75d image=nginx) k8s.ns=default k8s.pod=nginx
container=c73d9fc1a75d
```

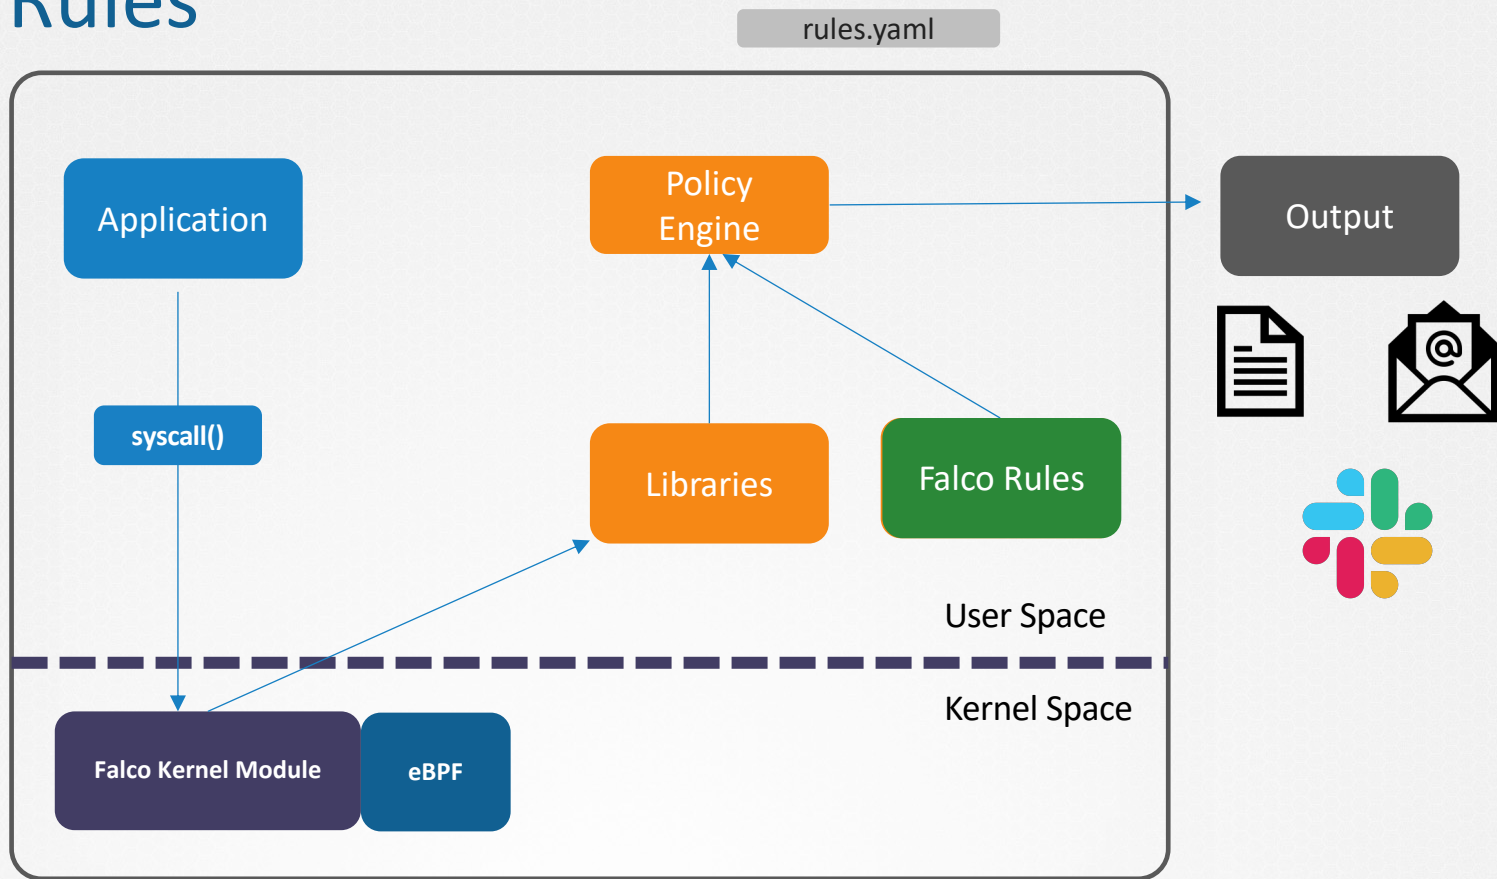
```
23:09:03.279503809: Warning Sensitive file opened for reading by non-trusted program (user=root
user_loginuid=-1 program=cat command=cat /etc/shadow file=/etc/shadow parent=bash gparent=runc
ggparent=containerd-shim gggparent=containerd-shim container_id=c73d9fc1a75d image=nginx)
k8s.ns=default k8s.pod=nginx container=c73d9fc1a75d k8s.ns=default k8s.pod=nginx
container=c73d9fc1a75d
```

Terminal 1

```
kubectl exec -ti nginx -- bash
```

```
root@nginx:/# cat /etc/shadow
```

Falco Rules



Falco Rules

rules.yaml

```
- rule: <Name of the Rule>
  desc: <Detailed Description of the Rule>
  condition: <When to filter events matching the rule>
  output: <Output to be generated for the Event>
  priority: <Severity of the event>
```


Falco Rules

rules.yaml

```
- rule: Detect Shell inside a container
  desc: Alert if a shell such as bash is open inside the container
  condition: container.id != host and proc.name = bash
  output: Bash Shell Opened (user=%user.name %container.id)
  priority: WARNING
```

Falco Rules

rules.yaml

```
- rule: Detect Shell inside a container
  desc: Alert if a shell such as bash is open inside the container
  condition: container.id != host and proc.name = bash
  output: Bash Opened (user=%user.name container=%container.id)
  priority: WARNING
```

container.id

proc.name

fd.name

evt.type

user.name

container.image.repository

rules.yaml

```
- rule: Detect Shell inside a container
  desc: Alert if a shell such as bash is open inside the container
  condition: {container.id!= host and proc.name = bash
  output: Bash Opened (user=%user.name container=%container.id)
  priority: {WARNING}
```

container.id

proc.name

fd.name

evt.type

username

container.image.repository

DEBUG

INFORMATIONAL

NOTICE


WARNING

ERROR

CRITICAL

ALERT

EMERGENCY


rules.yaml

```
- rule: Detect Shell inside a container
  desc: Alert if a shell such as bash is open inside the container
  condition: container.id != host and proc.name in (linux_shells)
  output: Bash Opened (user=%user.name container=%container.id)
  priority: WARNING
- list: linux_shells
  items: [bash, zsh, ksh, sh, csh]
```



rules.yaml

```
- rule: Detect Shell inside a container
  desc: Alert if a shell such as bash is open inside the container
  condition: container          and proc.name in (linux_shells)
  output:   Bash Opened (user=%user.name container=%container.id)
  priority: WARNING
- list: linux_shells
  items: [bash, zsh, ksh, sh, csh]
```

```
- macro: container
  condition: container.id != host
```




{K}ODE{K}LOUD

www.kodecloud.com

Falco Configuration Files



/etc/falco/falco.yaml

▶ journalctl -fu falco

```
-- Logs begin at Tue 2021-04-13 21:45:35 UTC, end at Tue 2021-04-13 21:51:31 UTC. --
Apr 13 21:45:36 node01 systemd[1]: Starting Falco: Container Native Runtime Security...
Apr 13 21:45:36 node01 systemd[1]: Started Falco: Container Native Runtime Security.
Apr 13 21:45:36 node01 falco[9817]: Falco version 0.28.0 (driver version 5c0b863ddade7a45568c0ac97d037422c9efb750)
Apr 13 21:45:36 node01 falco[9817]: Tue Apr 13 21:45:36 2021: Falco version 0.28.0 (driver version
5c0b863ddade7a45568c0ac97d037422c9efb750)
Apr 13 21:45:36 node01 falco[9817]: Falco initialized with configuration file /etc/falco/falco.yaml
Apr 13 21:45:36 node01 falco[9817]: Tue Apr 13 21:45:36 2021: Falco initialized with configuration file
/etc/falco/falco.yaml
```

/usr/lib/systemd/system/falco.service

```
[Unit]
Description=Falco: Container Native Runtime Security
Documentation=https://falco.org/docs/

[Service]
Type=simple
User=root
ExecStartPre=/sbin/modprobe falco
ExecStart=/usr/bin/falco --pidfile=/var/run/falco.pid -c /etc/falco/falco.yaml
.
```

```
/etc/falco/falco.yaml
```

```
#  
# Copyright (C) 2021 The Falco Authors.  
#  
#  
# Licensed under the Apache License, Version 2.0 (the "License");  
# you may not use this file except in compliance with the License.  
# You may obtain a copy of the License at  
#  
#   http://www.apache.org/licenses/LICENSE-2.0  
#  
.  
.  
.  
rules_file:  
  - /etc/falco/falco_rules.yaml  
  - /etc/falco/falco_rules.local.yaml  
  - /etc/falco/k8s_audit_rules.yaml  
  - /etc/falco/rules.d  
  
json_output: false  
log_stderr: true  
log_syslog: true  
log_level: info  
priority: debug
```

```
/etc/falco/falco.yaml
```

```
stdout_output:
  enabled: true

file_output:
  enabled: true
  filename: /opt/falco/events.txt

program_output:
  enabled: true
  program: "jq '{text: .output}' | curl -d @- -X POST https://hooks.slack.com/services/XXX"

http_output:
  enabled: true
  url: http://some.url/some/path/
```

```
/etc/falco/falco_rules.yaml
```

```
- rule: Terminal shell in container
  desc: A shell was used as the entrypoint/exec point into a container with an attached terminal.
  condition: >
    spawned_process and container
    and shell_procs and proc.tty != 0
    and container_entrypoint
    and not user_expected_terminal_shell_in_container_conditions
  output: >
    A shell was spawned in a container with an attached terminal (user=%user.name user_loginuid=%user.
loginuid %container.info
    shell=%proc.name parent=%proc.pname cmdline=%proc.cmdline terminal=%proc.tty container_id=%contain
er.id image=%container.image.repository)
  priority: NOTICE
.
.
.
```



```
/etc/falco/falco_rules.local.yaml
```

```
- rule: Terminal shell in container
  desc: A shell was used as the entrypoint/exec point into a container with an attached terminal.
  condition: >
    spawned_process and container
    and shell_procs and proc.tty != 0
    and container_entrypoint
    and not user_expected_terminal_shell_in_container_conditions
  output: >
    A shell was spawned in a container with an attached terminal (user=%user.name user_loginuid=%user.
loginuid %container.info
    shell=%proc.name parent=%proc.pname cmdline=%proc.cmdline terminal=%proc.tty container_id=%contain
er.id image=%container.image.repository)
  priority WARNING

- rule: Anomalous read in kodekloud/webapp pod
  desc: Detect Suspicious reads in custom webapp container
  condition: >
    open_read and container
    and container.image.repository="kodekloud/simple-webapp"
    and fd.directory != "/opt/app"
  output: >
    A file was opened and read outside the /opt/app directory(user=%user.name user_loginuid=%user.logi
nuid
    container_id=%container.id image=%container.image.repository)
  priority: CRITICAL
```

Hot Reload

```
▶ cat /var/run/falco.pid
```

```
7183
```

```
▶ kill -1 $(cat /var/run/falco.pid)
```

Hands-on Labs
cks.kodekloud.com



{K}ODE{K}LOUD

www.kodecloud.com

Kubernetes Auditing



```
kubectl logs -f falco-6t2dd
```

```
.  
.
.
```

```
22:57:09.163982780: Notice A shell was spawned in a container with an attached terminal (user=root  
user_loginuid=-1 k8s.ns=default k8s.pod=nginx container=c73d9fc1a75d shell=bash parent=runc  
cmdline=bash terminal=34816 container_id=c73d9fc1a75d image=nginx) k8s.ns=default k8s.pod=nginx  
container=c73d9fc1a75d
```

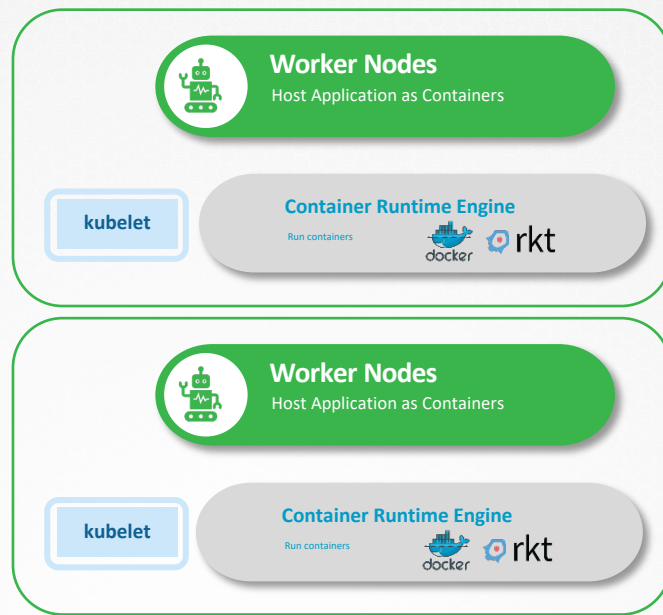

```
23:09:03.279503809: Warning Sensitive file opened for reading by non-trusted program (user=root  
user_loginuid=-1 program=cat command=cat /etc/shadow file=/etc/shadow parent=bash gparent=runc  
ggparent=containerd-shim gggparent=containerd-shim container_id=c73d9fc1a75d image=nginx)  
k8s.ns=default k8s.pod=nginx container=c73d9fc1a75d k8s.ns=default k8s.pod=nginx  
container=c73d9fc1a75d
```




namespace




```
kubectl run nginx --image nginx  
pod/nginx created
```



1. RequestReceived
2. RequestStarted
3. RequestComplete
4. Panic




webapp-pod deleted in prod-namespace!

audit-policy.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages:
rules:
```

1. RequestReceived




`webapp-pod deleted in prod-namespace!`

audit-policy.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages: ["RequestReceived"]
rules:
  - namespace: ["prod-namespace"]
```

1. RequestReceived




webapp-pod deleted in prod-namespace!

audit-policy.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages: ["RequestReceived"]
rules:
  - namespace: ["prod-namespace"]
    verb: ["delete"]
```

1. RequestReceived




webapp-pod deleted in prod-namespace!

audit-policy.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages: ["RequestReceived"]
rules:
  - namespace: ["prod-namespace"]
    verb: ["delete"]
    resources:
      - groups: " "
        resources: ["pods"]
```


1. RequestReceived




webapp-pod deleted in prod-namespace!

audit-policy.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages: ["RequestReceived"]
rules:
  - namespace: ["prod-namespace"]
    verb: ["delete"]
    resources:
      - groups: " "
        resources: ["pods"]
        resourceNames: ["webapp-pod"]
    level: RequestResponse
```


1. RequestReceived




webapp-pod deleted in prod-namespace!

audit-policy.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages: ["RequestReceived"]
rules:
  - namespace: ["prod-namespace"]
    verb: ["delete"]
    resources:
      - groups: " "
        resources: ["pods"]
        resourceNames: ["webapp-pod"]
    level: RequestResponse
  - level: Metadata
    resources:
      - groups: " "
        resources: ["secrets"]
```

1. RequestReceived



webapp-pod deleted in prod-namespace!

audit-policy.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages: ["RequestReceived"]
rules:
  - namespace: ["prod-namespace"]
    verb: ["delete"]
    resources:
      - groups: " "
        resources: ["pods"]
        resourceNames: ["webapp-pod"]
    level: RequestResponse
  - level: Metadata
    resources:
      - groups: " "
        resources: ["secrets"]
```



kube-apiserver.service

```
ExecStart=/usr/local/bin/kube-apiserver \\  
  --advertise-address=${INTERNAL_IP} \\  
  --allow-privileged=true \\  
  --apiserver-count=3 \\  
  --authorization-mode=Node,RBAC \\  
  --bind-address=0.0.0.0 \\  
  --enable-swagger-ui=true \\  
  --etcd-servers=https://127.0.0.1:2379 \\  
  --event-ttl=1h \\  
  --runtime-config=api/all \\  
  --service-cluster-ip-range=10.32.0.0/24 \\  
  --service-node-port-range=30000-32767 \\  
  --v=2  
  --audit-log-path=/var/log/k8-audit.log  
  --audit-policy-file=/etc/kubernetes/audit-policy.yaml
```

/etc/kubernetes/manifests/kube-apiserver.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  creationTimestamp: null  
  name: kube-apiserver  
  namespace: kube-system  
spec:  
  containers:  
    - command:  
      - kube-apiserver  
      - --authorization-mode=Node,RBAC  
      - --advertise-address=172.17.0.107  
      - --allow-privileged=true  
      - --enable-bootstrap-token-auth=true  
      - --audit-log-path=/var/log/k8-audit.log  
      - --audit-policy-file=/etc/kubernetes/audit-policy.yaml
```



kube-apiserver.service

```
ExecStart=/usr/local/bin/kube-apiserver \\  
  --advertise-address=${INTERNAL_IP} \\  
  --allow-privileged=true \\  
  --apiserver-count=3 \\  
  --authorization-mode=Node,RBAC \\  
  --bind-address=0.0.0.0 \\  
  --enable-swagger-ui=true \\  
  --etcd-servers=https://127.0.0.1:2379 \\  
  --event-ttl=1h \\  
  --runtime-config=api/all \\  
  --service-cluster-ip-range=10.32.0.0/24 \\  
  --service-node-port-range=30000-32767 \\  
  --v=2  
  --audit-log-path=/var/log/k8-audit.log  
  --audit-policy-file=/etc/kubernetes/audit-policy.yaml  
  --audit-log-maxage=10  
  --audit-log-maxbackup=5  
  --audit-log-maxsize=100
```

/etc/kubernetes/manifests/kube-apiserver.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  creationTimestamp: null  
  name: kube-apiserver  
  namespace: kube-system  
spec:  
  containers:  
    - command:  
      - kube-apiserver  
      - --authorization-mode=Node,RBAC  
      - --advertise-address=172.17.0.107  
      - --allow-privileged=true  
      - --enable-bootstrap-token-auth=true  
      - --audit-log-path=/var/log/k8-audit.log  
      - --audit-policy-file=/etc/kubernetes/audit-policy.yaml  
      - --audit-log-maxage=10  
      - --audit-log-maxbackup=5  
      - --audit-log-maxsize=100
```

audit-policy.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages:
  - "RequestReceived"
rules:
  - level: Metadata
    namespace: ["prod-namespace"]
    verb: ["delete"]
    resources:
      - group: ""
        resources: ["pods"]
```

```
{"kind":"Event","apiVersion":"audit.k8s.io/v1","level":"Metadata","auditID":"da2ad1a3-df15-4b10-a44d-79e73d7ec3c0","stage":"ResponseComplete","requestURI":"/api/v1/namespaces/prod-namespace/pods/webapp-pod","verb":"delete","user":{"username":"kubernetes-admin","groups":["system:masters","system:authenticated"]},"sourceIPs":["172.17.0.36"],"userAgent":"kubectl/v1.19.0 (linux/amd64) kubernetes/e199641","objectRef":{"resource":"pods","namespace":"prod-namespace","name":"webapp-pod","apiVersion":"v1"},"responseStatus":{"metadata":{},"code":200},"requestReceivedTimestamp":"2021-04-12T05:15:24.182178Z","stageTimestamp":"2021-04-12T05:15:24.182178Z"}
```



Hands-on Labs
cks.kodekloud.com



{K}ODE{K}LOUD

www.kodecloud.com

Immutable Infrastructure



Scripts



ANSIBLE

Mutable Infrastructure



Configuration Drift



v1.17



1

v1.17



2

v1.17



3







Immutable Infrastructure




Immutable Infrastructure

Dockerfile – My Custom Webapp

```
FROM nginx:1.19  
  
COPY nginx.conf /etc/nginx  
  
ENTRYPOINT ["sh", "entrypoint.sh"]
```

```
▶ kubectl cp nginx.conf nginx:/etc/nginx
```

Rolling Update





{K}ODE{K}LOUD

www.kodecloud.com

I

Ensure Immutability of Containers at Runtime



```
▶ kubectl cp nginx.conf nginx:/etc/nginx
```

```
▶ kubectl exec -ti nginx -- bash nginx:/etc/nginx  
root@nginx#
```

```
▶ kubectl create -f nginx.yaml
```

```
pod/nginx created
```

```
▶ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	Error	0	20s

nginx.yaml

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
    name: nginx
spec:
  containers:
    - image: nginx
      name: nginx
      securityContext:
        readOnlyRootFilesystem: true
```

```
▶ kubectl create -f nginx.yaml
```

```
pod/nginx created
```

```
▶ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	Error	0	20s

```
▶ kubectl logs nginx
```

```
root@controlplane:~# kubectl logs nginx
```

```
.  
.
```

```
2021/04/12 15:14:39 [emerg] 1#1: mkdir()  
"/var/cache/nginx/client_temp" failed (30: Read-only  
file system)
```

```
.
```

```
2021/04/12 16:11:26 [emerg] 1#1: open()  
"/var/run/nginx.pid" failed (30: Read-only file system)  
nginx: [emerg] open() "/var/run/nginx.pid" failed (30:  
Read-only file system)
```

nginx.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  labels:  
    run: nginx  
    name: nginx  
spec:  
  containers:  
  - image: nginx  
    name: nginx  
    securityContext:  
      readOnlyRootFilesystem: true
```

```
▶ kubectl create -f nginx.yaml
```

```
pod/nginx created
```

```
▶ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	Running	0	20s

```
▶ kubectl logs nginx
```

```
root@controlplane:~# kubectl logs nginx
.
.
2021/04/12 15:14:39 [emerg] 1#1: mkdir()
"/var/cache/nginx/client_temp" failed (30: Read-only
file system)
.
2021/04/12 16:11:26 [emerg] 1#1: open()
"/var/run/nginx.pid" failed (30: Read-only file system)
nginx: [emerg] open() "/var/run/nginx.pid" failed (30:
Read-only file system)
```

nginx.yaml

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
    name: nginx
spec:
  containers:
    - image: nginx
      name: nginx

  securityContext:
    readOnlyRootFilesystem: true

  volumeMounts:
    - name: cache-volume
      mountPath: /var/cache/nginx
    - name: runtime-volume
      mountPath: /var/run

  volumes:
    - name: cache-volume
      emptyDir: {}
    - name: runtime-volume
      emptyDir: {}
```



```
▶ kubectl cp nginx.conf nginx:/etc/nginx
```

```
tar: nginx.yaml: Cannot open: Read-only file system  
tar: Exiting with failure status due to previous errors  
command terminated with exit code 2
```

```
▶ kubectl exec -ti nginx - apt update
```

```
Reading package lists... Done  
E: List directory /var/lib/apt/lists/partial is missing. - Acquire (30: Read-only file system)  
command terminated with exit code 100
```

```
▶ kubectl create -f nginx.yaml
```

```
pod/nginx created
```

```
▶ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	20s

```
▶ kubectl exec -ti nginx - apt update
```

```
Reading package lists... Done
E: List directory /var/lib/apt/lists/partial is missing. - Acquire
(30: Read-only file system)
command terminated with exit code 100
```

nginx.yaml

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
  name: nginx
spec:
  containers:
    - image: nginx
      name: nginx

  securityContext:
    readOnlyRootFilesystem: true
    privileged: true
  volumeMounts:
    - name: cache-volume
      mountPath: /var/cache/nginx
    - name: runtime-volume
      mountPath: /var/run
  volumes:
    - name: cache-volume
      emptyDir: {}
    - name: runtime-volume
      emptyDir: {}
```

```
➤ kubectl create -f nginx.yaml
```

```
pod/nginx created
```

```
➤ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	20s

```
➤ kubectl exec -ti nginx - apt update
```

```
Reading package lists... Done
E: List directory /var/lib/apt/lists/partial is missing. - Acquire
(30: Read-only file system)
command terminated with exit code 100
```

```
➤ kubectl exec -ti nginx -- cat /proc/sys/vm/swappiness
```

```
60
```

```
➤ kubectl exec -ti nginx -- bash -c "echo '75' > /proc/sys/vm/swappiness"
```

```
➤ kubectl exec -ti nginx -- cat /proc/sys/vm/swappiness
```

```
75
```

```
node01
```

```
➤ cat /proc/sys/vm/swappiness
```

```
75
```

```
nginx.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
  name: nginx
spec:
  containers:
  - image: nginx
    name: nginx

  securityContext:
    readOnlyRootFilesystem: true
    privileged: true
  volumeMounts:
  - name: cache-volume
    mountPath: /var/cache/nginx
  - name: runtime-volume
    mountPath: /var/run

  volumes:
  - name: cache-volume
    emptyDir: {}
  - name: runtime-volume
    emptyDir: {}
```

readOnlyRootFilesystem: false



Privileged: true



runAsUser: 0



psp.yaml

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: example
spec:
  privileged: false
  readOnlyRootFilesystem: true
  runAsUser:
    rule: RunAsNonRoot
  seLinux:
    rule: RunAsAny
  supplementalGroups:
    rule: RunAsAny
  runAsUser:
    rule: RunAsNonRoot
  fsGroup:
    rule: RunAsAny
```

Hands-on Labs
cks.kodekloud.com



{K}ODE{K}LOUD

www.kodecloud.com